

# **CSE473: Computational Intelligence**

## **Bipolar Perceptron**

## **Learning & Optimization**

**by:**

**Hossam Abd El Munim**

**Computer & Systems Engineering Dept.,**

**Ain Shams University,**

**1 El-Sarayat Street, Abbassia, Cairo 11517**

# Dataset split

Training  
Images



Validation  
Images



Testing  
Images





- Train classifier

- Measure error
  - Tune model hyperparameters



- Secret labels
  - Measure error

*Random train/validate splits = cross validation*

# Training





# Testing





# Features


- Raw pixels
- Histograms
- Templates
- SIFT descriptors
  - GIST
  - ORB
  - HOG....





# Training




# Testing




# Training




# Testing



# Simulation of biological mechanism in artificial neural networks



(a) Biological neural network



(b) Artificial neural network

Figure 1.1: The synaptic connections between neurons. The image in (a) is from “*The Brain: Understanding Neurobiology Through the Study of Addiction* [598].” Copyright ©2000 by BSCS & Videodiscovery. All rights reserved. Used with permission.

# Deep Learning and Others





Figure 1.2: An illustrative comparison of the accuracy of a typical machine learning algorithm with that of a large neural network. Deep learners become more attractive than conventional methods primarily when sufficient data/computational power is available. Recent years have seen an increase in data availability and computational power, which has led to a “Cambrian explosion” in deep learning technology.


# **Humans Versus Computers: Stretching the Limits of Artificial Intelligence**

- Humans and computers are inherently suited to different types of tasks.
- For example, computing the cube root of a large number is very easy for a computer, but it is extremely difficult for humans.
- On the other hand, a task such as recognizing the objects in an image is a simple matter for a human, but has traditionally been very difficult for an automated learning algorithm.
- It is only in recent years that deep learning has shown an accuracy on some of these tasks that exceeds that of a human.
- In fact, the recent results by deep learning algorithms that surpass human performance [184] in (some narrow tasks on) image recognition would not have been considered likely by most computer vision experts as recently as 10 years ago.

# The Basic Architecture of Neural Networks



(a) Perceptron without bias



(b) Perceptron with bias

Figure 1.3: The basic architecture of the perceptron

# Perceptron as a Binary Classifier Only Two Categories




# Perceptron as a Binary Classifier (2D Case)

$$\mathbf{X} = [x_1 \ x_2 \ 1]^T$$


$$\mathbf{W} = [w_1 \ w_2 \ w_0]^T$$

$$\mathbf{W} \cdot \mathbf{X} = \mathbf{W}^T \mathbf{X} = 0$$

$$w_1 x_1 + w_2 x_2 + w_0 = 0$$



# Perceptron as a Binary Classifier (3D Case)



$$\mathbf{X} = [x_1 \ x_2 \ x_3 \ 1]^T$$

$$\mathbf{W} = [w_1 \ w_2 \ w_3 \ w_0]^T$$

$$\mathbf{W} \cdot \mathbf{X} = \mathbf{W}^T \mathbf{X} = 0$$

$$w_1x_1 + w_2x_2 + w_3x_3 + w_0 = 0$$

**What do you think about higher dimensions?**

# AND Function & Parameter Estimation

$$(-1)w_1 + (-1)w_2 + w_0 < 0$$

$$(-1)w_1 + (+1)w_2 + w_0 < 0$$

$$(+1)w_1 + (-1)w_2 + w_0 < 0$$

$$(+1)w_1 + (+1)w_2 + w_0 \geq 0$$


$x_1$	$x_2$	y
-1	-1	-1
-1	+1	-1
+1	-1	-1
+1	+1	+1

$$\text{Second} + \text{Third} = 2w_0 < 0$$

$$\text{First} + \text{Third} = -2w_2 + 2w_0 < 0$$

Take  $w_0 = -1$  & then  $w_2 > -1$ ...Set  $w_2 = 2$

Substitute in the first  $w_1 > -3$ ...Set  $w_1 = 2$



# Perceptron Training: Learning $\mathbf{W}$




Given a data set,  $D = \{(\mathbf{X}_1, y_1), (\mathbf{X}_2, y_2) \dots (\mathbf{X}_N, y_N)\}$  of labelled feature vectors where  $y \in \{-1, +1\}$ , we want to estimate the vector  $\mathbf{W}$  that minimizes the following objective function which measures the difference between the actual and desired outputs:-

$$E(\mathbf{W}) = (y_1 - \text{sign}(\mathbf{W} \cdot \mathbf{X}_1))^2 + (y_2 - \text{sign}(\mathbf{W} \cdot \mathbf{X}_2))^2 + \dots + (y_N - \text{sign}(\mathbf{W} \cdot \mathbf{X}_N))^2$$

# Training & Optimization

Given a data set,  $D = \{(\mathbf{X}_1, y_1), (\mathbf{X}_2, y_2) \dots (\mathbf{X}_N, y_N)\}$  of labelled feature vectors where  $y \in \{-1, +1\}$ , we want to estimate the vector  $\mathbf{W}$  that minimizes the following objective function which measures the difference between the actual and desired outputs:-

$$J(\mathbf{W}) = (y_1 - \text{sign}(\mathbf{W} \cdot \mathbf{X}_1))^2 + (y_2 - \text{sign}(\mathbf{W} \cdot \mathbf{X}_2))^2 + \dots + (y_N - \text{sign}(\mathbf{W} \cdot \mathbf{X}_N))^2$$




# Derivative-based Optimization

- We need to find  $W$  that satisfies the following condition:-

$$\frac{\partial J}{\partial W} = 0$$

- First, how can we solve such an equation in multi dimensions?
- Second, the sign function has a derivative of zero every where except at the origin. **Can we use a better objective function?**



# Again Learning W: An Optimization Problem

- Formulate learning problem as an optimization problems
  - Given: A set of  $N$  training examples
$$\{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_n, y_n)\}$$
  - Design a loss function  $L$
  - Find the weight vector  $\mathbf{w}$  that minimizes the expected/average loss on training data

$$J(\mathbf{w}) = \frac{1}{n} \sum_{m=1}^n L(\mathbf{w}^T \mathbf{x}_m, y_m)$$

# Better Loss Function

- 0/1 Loss function:  $J_{0/1}(\mathbf{w}) = \frac{1}{n} \sum_{m=1}^n L(\text{sign}(\mathbf{w}^T \mathbf{x}_m), y_m)$   
where  $L(y', y) = 0$  when  $y' = y$ , otherwise  $L(y', y) = 1$
- Issue: does not produce useful gradient since the surface of  $J$  is piece-wise flat
- Instead we will consider the “perceptron criterion”
$$J_p(\mathbf{w}) = \frac{1}{n} \sum_{m=1}^n \max(0, -y_m \mathbf{w}^T \mathbf{x}_m)$$
- The term  $\max(0, -y_m \mathbf{w}^T \mathbf{x}_m)$  is 0 when  $y_m$  is predicted correctly, otherwise it is equal to  $|\mathbf{w}^T \mathbf{x}_m|$ , which can be viewed as the “confidence” in the mis-prediction
- Has a nice gradient leading to the solution region

# Gradient Descent

- The objective function consists of a sum over data points--- we can update the parameter after observing each example
- This is the Stochastic gradient descent approach

$$J(\mathbf{w}) = \frac{1}{n} \sum_{m=1}^n \max(0, -y_m \mathbf{w}^T \mathbf{x}_m)$$
$$J_m(\mathbf{w}) = \max(0, -y_m \mathbf{w}^T \mathbf{x}_m)$$

$$\nabla J_m = \begin{cases} 0 & \text{if } y_m \mathbf{w} \cdot \mathbf{x}_m > 0 \\ -y_m \mathbf{x}_m & \text{otherwise} \end{cases}$$

After observing  $(\mathbf{x}_m, y_m)$ , if it is a mistake  $\mathbf{w} \leftarrow \mathbf{w} + y_m \mathbf{x}_m$

# Batch Training

Input:  $(\mathbf{X}_m, y_m)$ ,  $m = 1, 2, \dots, N$

Set  $\mathbf{W} = [0, 0, \dots, 0]^T$

Repeat

**delta** =  $[0, 0, \dots, 0]^T$

for  $m = 1$  to  $N$  do

    if  $y_m \mathbf{W} \cdot \mathbf{X}_m \leq 0$

**delta** = **delta** –  $y_m \mathbf{X}_m$

**delta** = **delta** /  $N$

**W** = **W** – **delta**

Until  $\|\mathbf{delta}\| < \epsilon$

# Online Training

Input:  $(\mathbf{X}_m, y_m)$ ,  $m = 1, 2, \dots, N$

Set  $\mathbf{W} = [0, 0, \dots, 0]^T$

Repeat

**delta** =  $[0, 0, \dots, 0]^T$

for  $m = 1$  to  $N$  do

    if  $y_m \mathbf{W} \cdot \mathbf{X}_m \leq 0$

**delta** = **delta** –  $y_m \mathbf{X}_m$


$\mathbf{W} = \mathbf{W} - \mathbf{delta} / N$

Until  $\|\mathbf{delta}\| < \epsilon$

# Bipolar Perceptron and SVM

Find a *linear function* to separate the classes:

$$f(\mathbf{x}) = \text{sgn}(\mathbf{w} \cdot \mathbf{x} + b)$$




What if my data are not linearly separable?

Introduce flexible ‘hinge’ loss (or ‘soft-margin’)

# Relation to SVM: Modified Loss Function

$$J(\mathbf{w}) = \frac{1}{n} \sum_{m=1}^n \max(0, 1 - y_m \mathbf{w}^T \mathbf{x}_m)$$



Introduce flexible ‘hinge’ loss (or ‘soft-margin’)

# Relation to SVM: Modified Loss Function

Does this affect the training algorithm?

# Choice of Activation Function (1)




Figure 1.7: Pre-activation and post-activation values within a neuron


# Choice of Activation Function (2)

The classical activation functions that were used early in the development of neural networks were the sign, sigmoid, and the hyperbolic tangent functions:


$$\Phi(v) = \text{sign}(v) \text{ (sign function)}$$

$$\Phi(v) = \frac{1}{1 + e^{-v}} \text{ (sigmoid function)}$$


$$\Phi(v) = \frac{e^{2v} - 1}{e^{2v} + 1} \text{ (tanh function)}$$



(a) Identity



(b) Sign



(c) Sigmoid


# Choice of Activation Function (3)

The sigmoid and the tanh functions have been the historical tools of choice for incorporating nonlinearity in the neural network. In recent years, however, a number of piecewise linear activation functions have become more popular:


$$\Phi(v) = \max\{v, 0\} \text{ (Rectified Linear Unit [ReLU])}$$

$$\Phi(v) = \max \{\min [v, 1], -1\} \text{ (hard tanh)}$$


The ReLU and hard tanh activation functions have largely replaced the sigmoid and soft tanh activation functions in modern neural networks because of the ease in training multi-layered neural networks with these activation functions.



(d) Tanh



(e) ReLU



(f) Hard Tanh

Figure 1.8: Various activation functions

# Sigmoid and tanh Derivatives

Sigmoid ( $u$ ) =  $f(u)$

$df/du = f(1-f)$


tanh ( $v$ ) =  $g(v)$

$dg/dv = (1-g^2)$

# Will it work? Do need a multilayer network?




LINEARLY SEPARABLE



NOT LINEARLY SEPARABLE

# A Multilayer Network



# Multiple Outputs




Figure 1.9: An example of multiple outputs for categorical classification with the use of a softmax layer

# Probabilistic Prediction


## Index of Maximum is the Prediction

- Softmax layer as the output layer

**Probability:**

- $1 > y_i > 0$
- $\sum_i y_i = 1$

### Softmax Layer



# Choice of Loss Function for Probabilistic Prediction

# Two Categories: Binary Classification Identity Activation Function

**Binary targets (logistic regression):** In this case, it is assumed that the observed value  $y$  is drawn from  $\{-1, +1\}$ , and the prediction  $\hat{y}$  is a an arbitrary numerical value on using the identity activation function. In such a case, the loss function for a single instance with observed value  $y$  and real-valued prediction  $\hat{y}$  (with identity activation) is defined as follows:

$$L = \log(1 + \exp(-y \cdot \hat{y})) \quad (1.14)$$


# Two Categories: Binary Classification Sigmoid Activation Function

This type of loss function implements a fundamental machine learning method, referred to as *logistic regression*. Alternatively, one can use a sigmoid activation function to output  $\hat{y} \in (0, 1)$ , which indicates the probability that the observed value  $y$  is 1. Then, the negative logarithm of  $|y/2 - 0.5 + \hat{y}|$  provides the loss, assuming that  $y$  is coded from  $\{-1, 1\}$ . This is because  $|y/2 - 0.5 + \hat{y}|$  indicates the probability that the prediction is correct. This observation illustrates that one can use various combinations of activation and loss functions to achieve the same result.


$$L = -\log\left(\left|\frac{y}{2} - \frac{1}{2} + \hat{y}\right|\right)$$

$$\hat{y} = \frac{1}{1 + e^{-(\mathbf{W} \cdot \mathbf{X})}}$$

# Two Categories: Binary Classification Sigmoid Activation Function ( $y=1$ )



# Two Categories: Binary Classification Sigmoid Activation Function ( $y = -1$ )



# Multi-Classifications

**Categorical targets:** In this case, if  $\hat{y}_1 \dots \hat{y}_k$  are the probabilities of the  $k$  classes (using the softmax activation of Equation 1.9), and the  $r$ th class is the ground-truth class, then the loss function for a single instance is defined as follows:

$$L = -\log(\hat{y}_r) \tag{1.15}$$

# Machine Learning with Shallow Neural Networks

# Neural Architectures for Binary Classification Models






Figure 2.2: An extended architecture of the perceptron with both discrete and continuous predictions

# Different Variants of the Perceptron

## Binary Classifier – Bad Loss




# Different Variants of the Perceptron Binary Classifier – Smooth Loss Function




# Different Variants of the Perceptron

## Linear Regression – SSD Loss Function




# Different Variants of the Perceptron

## Logistic Regression – LOG Likelihood Loss Function



# Different Variants of the Perceptron

## Logistic Regression – LOG Likelihood Loss Function (Alternative)



# Different Variants of the Perceptron Binary Classifier – Hinge Loss (SVM)

